

Image Coordinate Systems in AIPS++

Mark Calabretta

AIPS++ programmer group
1992/09/16

1 Introduction

This document discusses the design and implementation of coordinate systems for images in AIPS++. Loosely speaking, an AIPS++ “image” has a regularity which allows that coordinates do not need to be stored with each pixel. In this sense, images differ from the more general data structure, such as that used for uv-data, for which coordinates must be stored for every data element.

Another property of an AIPS++ image is that of interpolability - adjacent pixels are related to one another in such a way that it makes sense to interpolate a value between them. This is not a property possessed by “IF”, “polarization”, or “source id” axes, which may instead be regarded as multi-valued pixels (polarization), or lists of images (IF, and source id). These may produce regular data structures which are superficially similar to an image, but some part of the “coordinate” information would be stored in the descriptor for the pixel values or list elements as the case may be.

In a simple view of a coordinate system, part of an image’s regularity may be seen as arising from the data elements, or pixels, being separated by integral multiples of some fundamental increment in each of its N axes. This is the model adopted by FITS; each axis of an N-dimensional image has a coordinate reference pixel, the coordinate value at that pixel, and a coordinate increment. An additional parameter may also be used to describe a rotation of coordinates although this is not rigorously defined. In this model, non-orthogonal pairs of coordinate axes, such as right ascension and declination in a polar projection, are treated as special instances of the simple case.

However, the FITS view of coordinate systems sits uncomfortably with non-linear coordinates in 1-dimension, such as might arise from frequency-velocity axes, logarithmic axes, and so on. The difficulties are compounded by pairs, triples, or higher aggregations of dependent axes, as might arise from most conventional map projections of celestial coordinates. In fact, the FITS methodology cannot handle a general linear transformation of a set of “orthogonal” linear axes.

It is worth discussing the various shades of intergradation between what is considered to be a regular image, and the general data structure characterized by uv data.

- At the simplest level we may have a collection of pixels which, by virtue of the way they are organized, and the nature of the coordinate system itself, does not need coordinates to be stored with each pixel.

It is possible for a regular N-dimensional array of numbers to have this property, and this forms the conventional view of a regular image. Since the individual pixels are addressable by *pixel coordinates*, that is, integral array indices (p_1, p_2, \dots) , the only other requirement is a pre-defined algorithm for converting these to *image coordinates* (i_1, i_2, \dots) .

Run-length lists also allow integral indexing, although via a somewhat more complicated mechanism. Instances of these can also fall into this simple category.

- Consideration of what actually constitutes the algorithm for computing $\mathbf{I} = (i_1, i_2, \dots)$ from $\mathbf{P} = (p_1, p_2, \dots)$ leads us to a higher level of abstraction. Let us write the *mapping function*, $f()$, as

$$\mathbf{I} = f(\mathbf{P}) \quad (1)$$

where \mathbf{I} , and \mathbf{P} are vector quantities. In the simple case where $f()$ is completely separable we may write

$$(i_1, i_2, \dots) = (f_1(p_1), f_2(p_2), \dots). \quad (2)$$

This is the mathematical condition for what is usually referred to as “orthogonal” or “rectilinear” coordinates. However, the term “separable” is more mathematically meaningful. Separability is at the heart of the FITS coordinate model, celestial coordinates and linear transformations being an addition to it. FITS also assumes that the functions $f_1()$, $f_2()$, etc. can be expressed as linear equations involving a few coefficients, although presumably it could be extended to include the notion of logarithmic, inverse, or other non-linear coordinate axes.

However, suppose that some or all of the functions $f_1()$, $f_2()$,... could not be expressed in mathematical form, but instead could only be defined in terms of a table lookup. In terms of an AIPS++ image it is still not necessary to store image coordinates with every pixel, although rather more storage is required to define the mapping function.

FITS currently solves this problem by the use of *random parameters*.

- Taking the previous example to its logical conclusion we must consider the case where the mapping function is not separable, and is expressible only as a lookup table. In the completely degenerate case where table entries are required for each pixel, we arrive at the most general structure where coordinates are stored with each pixel. (It may, however, still be meaningful to talk of such an image as being regular if, for example, adjacent pixels are not independent of each other, that is, if it makes sense to interpolate them. However, this is a fine point which is outside the scope of the present discussion.)

Between these extremes there exists a wide range of possibilities in which the mapping function may be partly separable, and need be expressible only partly via lookup tables.

2 Requirements

The following examples illustrate how coordinate systems might be used in AIPS++.

- Determination of image coordinates. For example, “If the peak emission is at pixel coordinate (35.74, 129.05, 11.10), what are the corresponding image coordinates?”
Answer: (150.33956, -35.49803, 847.0145, 8.5729×10^{-4}), note more image than pixel coordinate elements.
- Description of image coordinates - their type, and units. For example, “How do I interpret the image coordinates returned in the above example?”
Answer: ecliptic longitude (degrees), ecliptic latitude (degrees), heliocentric velocity (km/s), and LSR velocity (pc/yr).

In this example where the latter pair of coordinates are dependent, a mechanism will be required to indicate which of them is to be considered “independent”. Alternatively, in cases like that of (l, m, n) direction cosines where one of the triple is dependent on the other two ($l^2 + m^2 + n^2 = 1$), it may be more convenient to indicate which of them is to be considered “dependent”.

Queries could also be of the form, “Which axes of my 7-dimensional image correspond to right ascension and declination?”, or more generally, “Which axes of my 7-dimensional image correspond to the sky plane?”.

- Determination of pixel coordinates from image coordinates. For example, plotting the positions of a list of sources on a map.
- Plotting of coordinate grids. This is an extension of determining pixel coordinates but may also involve using the mapping function in a different way. For example, to determine the point where a grid line of constant galactic latitude intersects the frame defined as a line (or in general a surface) where one of the pixel coordinate elements is constant. The point of intersection is defined by a mix of image and pixel coordinate elements.
- Transformation of coordinates. For example, the mapping function for a spectral-line cube which returns right ascension, declination, and frequency could be modified so that it returns galactic longitude, galactic latitude, and barycentric velocity. Precession of equatorial coordinates is another example, as is changing the units of the returned coordinate elements (e.g. degrees \leftrightarrow radians, kilometers \leftrightarrow miles).
- General linear transformations. For example, translating, scaling, or rotating the coordinate axes (as opposed to regridding the pixels).
- Interpolation of an image onto a new grid so that its coordinate system transforms to a specified coordinate system, in particular, that of another image. This will be required to overlay images with different coordinate systems.
- Extraction of an interpolated “flat” n-dimensional subimage of arbitrary orientation from an m-dimensional image, for example, an oblique 2-D plane from a 3-D data cube. Although interpolation would be outside the function of the coordinate system class itself, it must be able to handle the new coordinate system created by this operation. The extraction of a single interpolated pixel value may be considered a special case of this requirement.

This example shows how degenerate mapping functions may arise, since the association of each pixel in the oblique n-dimensional subimage with its m-dimensional image coordinate may need to be preserved.

- Extraction of a “curved” n-dimensional subimage from an m-dimensional image. Since a curved surface cannot, in general, be “flattened” out without distortion, it will be necessary to take the m-dimensional subimage which just contains the surface and define the curved surface as a subset of this (via a function or lookup table).

There are, however, some notable special cases. For instance a curved 1-dimensional subimage can always be “straightened out” and stored in a 1-dimensional vector. Likewise, a corrugated 2-dimensional subimage can be flattened into a 2-dimensional array. However, since the coordinates in these instances would almost certainly have to be stored in a look-up table, it may be preferable to preserve the embedding m-dimensional subimage even here.

- FITS input and output must be supported.
- Ancillary functions should be provided for reading and reporting coordinate values in the required form and precision. For example, formatting angles in DD:MM:SS.S format to the required number of decimal places, or, as for a right ascension, in HH:MM:SS.SSS time format, or as a decimal hour, or decimal degrees. Negative angles with a degree field of zero must be handled correctly, for example $-00^{\circ}05'27''.554$. Dates and times should be reported in the standard AIPS++ date format.

It should be possible to report coordinates in the units specified. For example, a coordinate element returned by the mapping function in light years could be reported in parsec. This includes controlling the metric prefix (where relevant), for example, m (meters) or km (kilometers), pc (parsec) or Mpc (Megaparsec).

Sufficient information should be stored with the image coordinates to allow sensible default formatting.

3 Analysis

The operations described above would not necessarily all be included within a coordinate class, `CoordSys`, but might be implemented as methods of an image class, `Image`, which uses `CoordSys`. For example, the operation of regridding an image to conform with a predefined coordinate system does not belong within `CoordSys` but would rely heavily on the methods provided by it. In this section we will attempt to further refine our understanding of what the `CoordSys` class consists of.

3.1 Spherical Coordinate Systems

Coordinate systems which represent the celestial sphere are of special significance in astronomy and provide a good illustration of an important point concerning the coordinate mapping function. Such coordinate systems consist of a spherical coordinate system together with a (spherical) map projection. The two components are separate entities, and there are good reasons to treat them as such.

Although map projections can only be specified mathematically with reference to a spherical coordinate system, they are in fact geometrical entities which exist independently of any coordinate system. This is most clearly seen in the “projective” map projections which may be defined by purely geometrical constructs. The particular spherical coordinate system chosen when representing them mathematically has no special significance, since the projection could be reformulated in terms of any other spherical coordinate system.

Practically speaking, changing an image’s map projection, for example changing its obliquity, or changing its type from say orthographic (SIN) to gnomonic (TAN), always requires interpolation on the pixel values and usually changes the shape of objects in the map. On the other hand, transforming an image’s coordinate system, for example from equatorial to ecliptic, does not require interpolation and does not distort objects but instead simply introduces a new coordinate grid.

In fact, the dichotomy between the “static” and “transformable” parts of the mapping function can be seen in other places. For example, the wavelength calibration of an optical spectrum would be the static part of the spectrum’s mapping function, whereas expression of the coordinate as a wavelength, frequency, or velocity, or changing its physical units (\AA or nm), would be the transformable part. The static part of the mapping function will henceforth be referred to as the *coordinate structure function*, and the transformable part as the *coordinate transformation function*.

3.2 Distillation of requirements

At a fairly abstract level, AIPS++ coordinate systems must have the following features:

- The coordinate system should be definable via the mapping function $f()$, via its inverse, or both (for processing efficiency). It must be possible to build the coordinate system (i.e. define $f()$ or its inverse) at run time in response to external requests such as FITS input, image manipulation, user activity, etc.

Degenerate mappings such as might correspond to $(RA, dec) \rightarrow (l, m, n)$, and overspecified mappings such as $(l, m, n) \rightarrow (RA, dec)$ must be provided for.

- The following functions for transforming coordinates will be needed within `CoordSys`:

- *forward mapping function*, $f()$, which takes \mathbf{P} and returns \mathbf{I} ,

- *inverse mapping function*, $f^{-1}()$, which takes \mathbf{I} and returns \mathbf{P} , and
- *mixed mapping function*, $\varphi()$, which when given some of the (p_1, p_2, \dots) and some of the (i_1, i_2, \dots) returns the others.

Although the mixed mapping function encompasses both the mapping and inverse mapping function, it may well be based upon them.

- Transformation of coordinate systems should be done within the **CoordSys** class. For example, if a mapping function $f_e(\mathbf{P})$ returns vector \mathbf{I}_e in equatorial coordinates, it should be possible to tell the **CoordSys** object to transform $f_e()$ to $f_g()$ which returns vector \mathbf{I}_g in galactic coordinates.
- It will often be required to transform coordinates which exist without reference to any image. For example, transforming a list of source positions in equatorial coordinates to galactic coordinates does not relate to any image, but is intimately related to the coordinate transformation function part of the mapping function. The transformation function should therefore be accessible outside the context of **CoordSys** for use by **Coordinate** objects for general coordinate transformations.

Coordinate iterators should be supplied for the **Coordinate** class.

- General linear transformations should be handled within **CoordSys**. Translations can be handled by adding an extra dimension to the transformation matrix, as is done with the “CTM” (current transformation matrix) in POSTSCRIPT.
- Subimage coordinate systems, including those of oblique subimages, can be derived from the parent image by transforming the linear transformation matrix which forms a part of the mapping function (see the following example).
- Where the dimensionality of a subimage is less than that of its parent image, the **CoordSys** object of the subimage may have to be derived from the parent but with dimensional restrictions. For example, an oblique line (slice) extracted from a 2-dimensional image has only one axis, but clearly it might be sensible to associate the coordinates of the two axes of the parent image (say right ascension and declination) with each point on the slice axis.

One way to do this would be to copy the parent image’s **CoordSys** object to the subimage, apply a rotation, translation, and scale (to account for the interpolation interval) via the linear transformation matrix so that the oblique slice was mapped into the x -pixel axis of the parent image, and then flag the y -pixel coordinate of the subimage as having a constant value of zero (the restriction).

Changes in the relative order of the subimage axes, i.e. transpositions, could be implemented via transpositions of the rows and columns of the transformation matrix.

- Since it knows all about the image coordinate system, the **CoordSys** class should respond to queries concerning the coordinate axis types and their units.
- FITS output should be supported on a best effort basis. **CoordSys** will need a method which can determine whether the mapping function is simple enough to be cast into FITS header form and if so return the associated header cards.

If the coordinate system is too complex for the FITS coordinate model, the AIPS++ FITS writer will have to resort to using random parameters.

- Processing efficiency considerations: although the **CoordSys** class will be involved in some difficult numerical computations it is unlikely that these will be so numerous as to cause a processing bottleneck in themselves.

However, certain mathematical transformations such as FFTs and DFTs are based implicitly on linear coordinate systems and are most efficiently expressed in terms of the coefficients of these linear systems (reference pixel, reference value, and increment for each axis). The alternative of computing the coordinates of each pixel within the inner processing loop of these algorithms is far too inefficient.

Therefore, the `CoordSys` class must be prepared to supply the coefficients of these linear coordinate systems if it is sensible to do so (this is closely related to the task of producing FITS headers).

3.3 Storage coordinates

The relationship between pixel coordinates and the actual storage location of the pixels was not discussed in the above analysis which concerned itself only with the relationship between pixel and image coordinates. Put another way, no assumption has been made concerning the value of the pixel coordinate of the first pixel in the image (top-, or bottom-left hand corner, depending on which is adopted for AIPS++).

It may be useful for the `Image` class to differentiate between pixel coordinates and *storage coordinates* which are related to the way pixels are stored and so are inherently integral. For example, the first pixel in an image would always have storage coordinates (1, 1, ...). Pixel coordinates might be translated and possibly inverted with respect to storage coordinates.

Some examples of why it may be useful to distinguish between pixel and storage coordinates are:

- Where an image was composed of a mosaic of subimages it would be easy to define coordinate systems for all subimages by setting their pixel coordinates to correspond to the pixel coordinates of the larger canvas and copying the `CoordSys` object for each of them from the canvas (or pointing them to it).
- Conversely, testing for equality of coordinate systems of two images (perhaps non-overlapping elements of a mosaic) would revert to testing for equality of their `CoordSys` objects.
- Constructing the coordinate system for a non-oblique subimage would be as easy as resetting the offsets between storage and pixel coordinates and copying the `CoordSys` object from the parent image.
- Translated coordinate systems often arise in image display, especially where panning, scrolling, or split-screen is involved. In this case storage coordinates might correspond to TV pixel coordinates.

Storage coordinates are outside the scope of the `CoordSys` class which deals only with the relationship between pixel and image coordinates, and it would be the responsibility of the `Image` class to implement them. This would entail the maintenance of the integral offsets between pixel coordinates and storage coordinates, plus a query function to report the pixel coordinates at each corner of the image. Applications programmers would need to bear in mind that the pixel coordinates of an image don't necessarily begin at (1, 1, ...), and with each request for the pixel values from a region of the image they would be supplied with the corresponding pixel coordinates. To save confusion, storage coordinates should only appear in low-level `Image` methods (possibly as private data) and remain invisible to applications programmers.

3.4 Object model

A Rumbaugh object model diagram for the AIPS++ coordinate classes is presented in fig 1. The mapping function, which forms the heart of the diagram, is divided into an ordered sequence of three components, a *linear transformation*, followed by the *coordinate structure function*, and then the *coordinate transformation function*. The *subimaging* association is modelled as the *linear transformation* class, and likewise the *coordinate transformation* association is modelled as the *coordinate transformation function* class. The relation "mapping function transforms pixel coordinate to image coordinate" is represented as a ternary association between these three classes. Also of note, the association between `CoordSys` and `Image` is one-to-many to allow for a `CoordSys` object to be shared amongst more than one `Image` object, for example, a dirty map, dirty beam, and cleaned maps.

Figure 1: Rumbaugh object model diagram for the AIPS++ coordinate classes.

3.5 Dynamic model

The dynamic model for the AIPS++ coordinate classes is trivial since they implement a non-interactive computation (the mapping function).

3.6 Functional model

A Rumbaugh functional model diagram for the AIPS++ coordinate classes is presented in fig 2. It provides a coarse-grained description of the operation of the mapping function, and shows the operation of each of its three components.

Figure 2: Rumbaugh functional model diagram for the AIPS++ coordinate classes.

3.7 Data dictionary

Terms used in this document are listed here in glossary form.

- *Coordinate system* - an association between a set of points and a set of tuples of numbers.

- *Coordinate structure function* - the static part of the *mapping function* representing the physical distortion of an image, for example a spherical map projection or spectral wavelength calibration. It can only be changed by interpolating the image onto a new grid.
- *Coordinate transformation* - a change in the association between the set of points and the set of tuples of numbers which comprise a *coordinate system*.
- *Coordinate transformation function* - the transformable part of the *mapping function*. It implements a *coordinate transformation*.
- *CoordSys* - class of objects containing information which completely describes image coordinate systems, including the axis types (frequency, right ascension, etc.), their physical units (MHz, radians, etc.), and the algorithm for transforming between *pixel coordinates* and *image coordinates*.
- *Image coordinate* - a tuple of real numbers which defines a point in a coordinate system associated with an image. The image coordinate system usually represents measurable physical parameters such as time, distance, and angle, and the coordinate elements are the measured values of these parameters with associated units.
- *Inverse mapping function* - see *Mapping function*.
- *Linear transformation* - any operation which can be represented by a matrix with constant elements. This includes translation, scaling, rotation, and axis skew. The *mapping function* incorporates a linear transformation matrix which, in particular, may be used to define the coordinate system of a subimage in terms of its parent image.
- *Mapping function* - the algorithm used to transform *pixel coordinates* to *image coordinates*. The *inverse mapping function* transforms *image coordinates* to *pixel coordinates*, and the *mixed mapping function* is a hybrid of the two.
- *Mixed mapping function* - see *Mapping function*.
- *Pixel coordinate* - a tuple of real numbers which locates a pixel in an image. The pixel coordinate elements of a pixel are integral, and fractional coordinate elements span the interval between pixels. Pixel coordinates may have an integral offset from *storage coordinates*.
- *Storage coordinate* - a tuple of integers which locates a pixel in an image. The first pixel in an image has storage coordinate (1, 1, ...).